

# Contemporary Computer Architecture TDSN13

---

LECTURE 4 – INTRODUCTION TO CUDA

ANDREAS AXELSSON (ANDREAS.AXELSSON@JU.SE)

# Introduction to CUDA


---

- CUDA Architecture
  - Expose GPU parallelism for general-purpose computing
  - Retain performance
- CUDA C/C++
  - Based on industry-standard C/C++
  - Small set of extensions to enable heterogeneous programming
  - Straightforward APIs to manage devices, memory etc.
- How to get started

# Heterogeneous Computing

---

- Terminology:
  - *Host* The CPU and its memory (host memory)
  - *Device* The GPU and its memory (device memory)



Host



Device

# CUDA Parallel Computing Platform

Programming  
Approaches

Libraries

“Drop-in”  
Acceleration


OpenACC  
Directives

Easily Accelerate  
Apps

Programming  
Languages

Maximum Flexibility

Development  
Environment



Nsight IDE  
Linux, Mac and Windows  
GPU Debugging and  
Profiling

CUDA-GDB  
debugger  
NVIDIA Visual  
Profiler

Open Compiler  
Tool Chain




Enables compiling new languages to CUDA  
platform, and CUDA languages to other  
architectures

Hardware  
Capabilities


SMX




Dynamic  
Parallelism



HyperQ



GPUDirect



# CUDA GPU Programming

---

**Fortran** ▶

CUDA Fortran

**C** ▶

CUDA C

**C++** ▶

CUDA C++

**Python** ▶

PyCUDA

[Home](#) > [High Performance Computing](#) > [CUDA Toolkit](#) > [CUDA Toolkit 10.2 Download](#)

## Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown.


Operating System

[Windows](#)[Linux](#)[Mac OSX](#)[Documentation >](#)[Release Notes >](#)[Code Samples >](#)[Legacy Releases >](#)


## Get Started

The above options provide the complete CUDA Toolkit for application development. Runtime components for deploying CUDA-based applications are available in [ready-to-use containers from NGC](#).

### Installing the CUDA Toolkit



### Introduction to CUDA



### Getting Started with CUDA




### Discover Latest CUDA Capabilities



# Nvidia – GPU Architectures

---



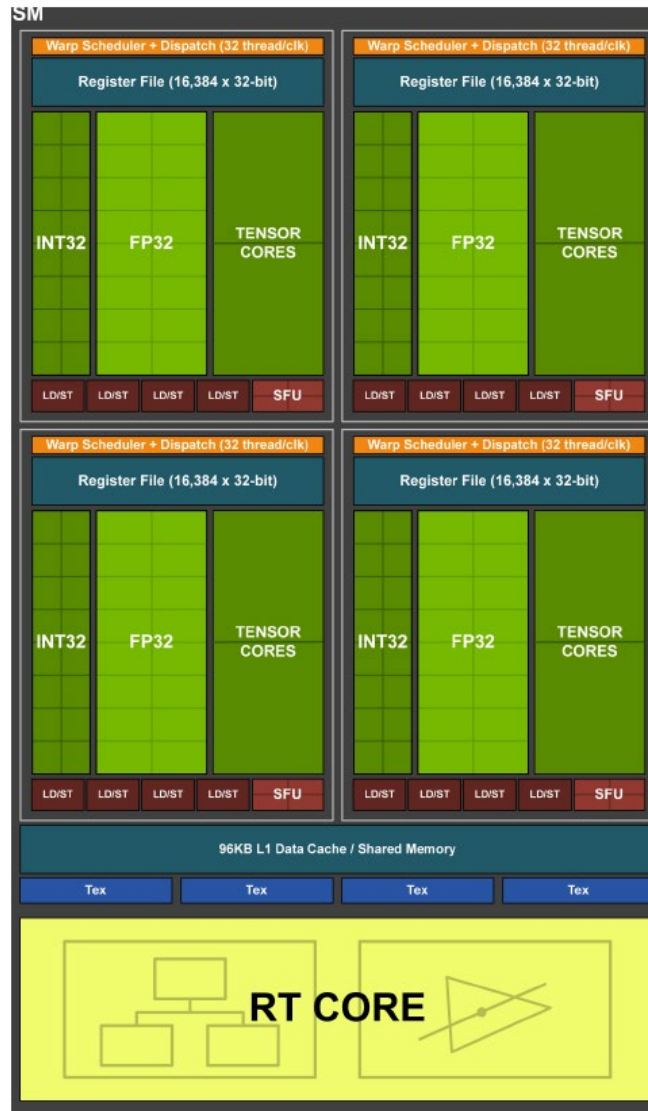
- Tesla
- Fermi
- Kepler
- Maxwell
- Pascal
- Volta
- Turing
- Ampere?

# Turing Architecture






# Turing Streaming Multiprocessor



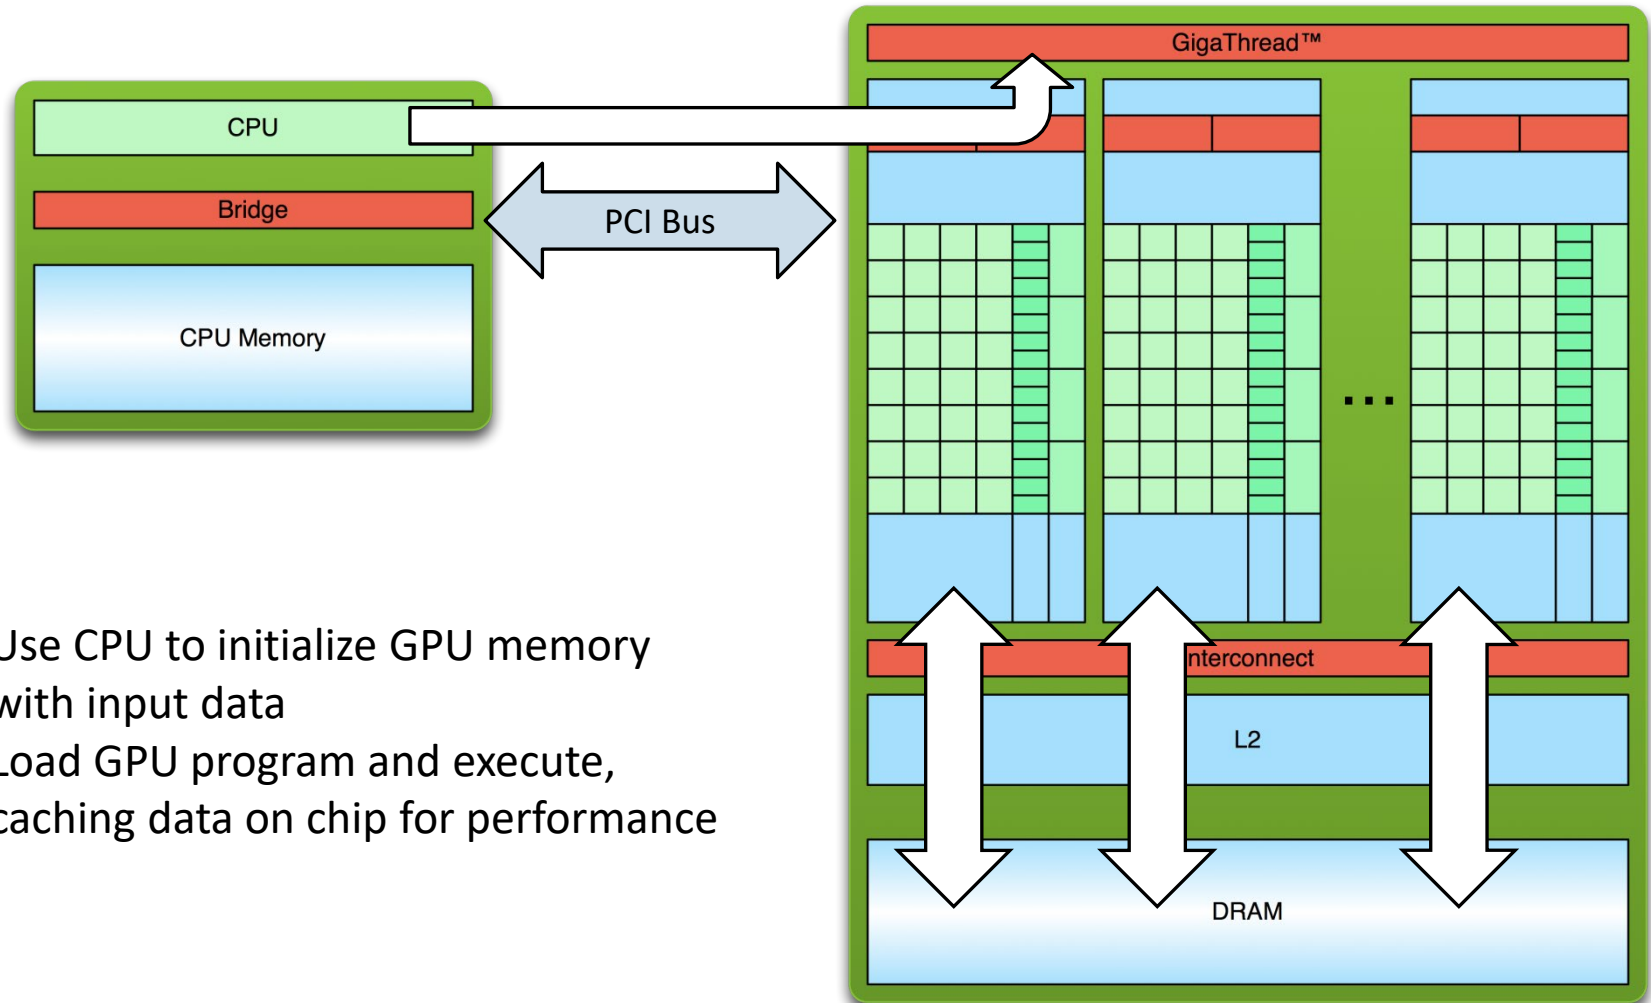
# Workflow

---

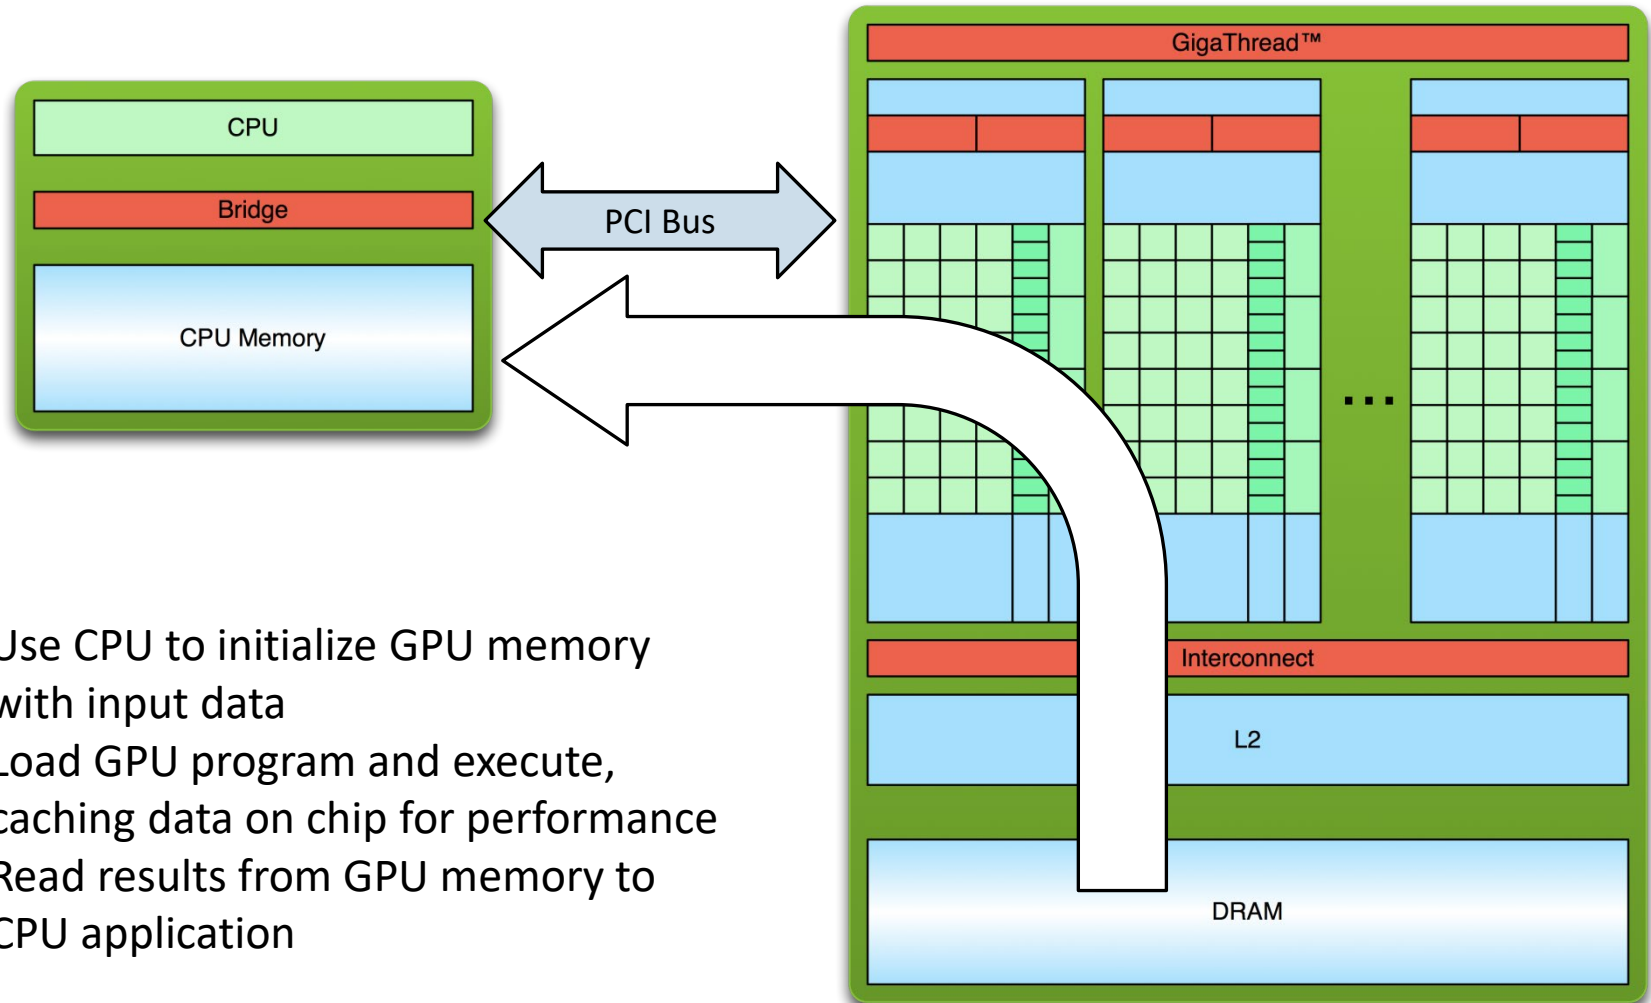


1. Use CPU to initialize GPU memory with input data

# Workflow



# Workflow



# Hello World!

---

```
int main(void) {  
    printf("Hello World!\n");  
    return 0;  
}
```

- Standard C that runs on the host
- NVIDIA compiler (nvcc) can be used to compile programs with no *device* code

Output:

```
$ nvcc  
hello_world.  
cu  
$ a.out  
Hello World!  
$
```

# Hello World! with Device Code

---

```
__global__ void mykernel(void) {  
}
```

```
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

- Two new syntactic elements...

# Hello World! with Device Code

---

```
mykernel<<<1,1>>>();
```

- Triple angle brackets mark a call from *host* code to *device* code
  - Also called a “kernel launch”
  - We’ll return to the parameters (1,1) in a moment
- That’s all that is required to execute a function on the GPU!

# Hello World! with Device Code

---

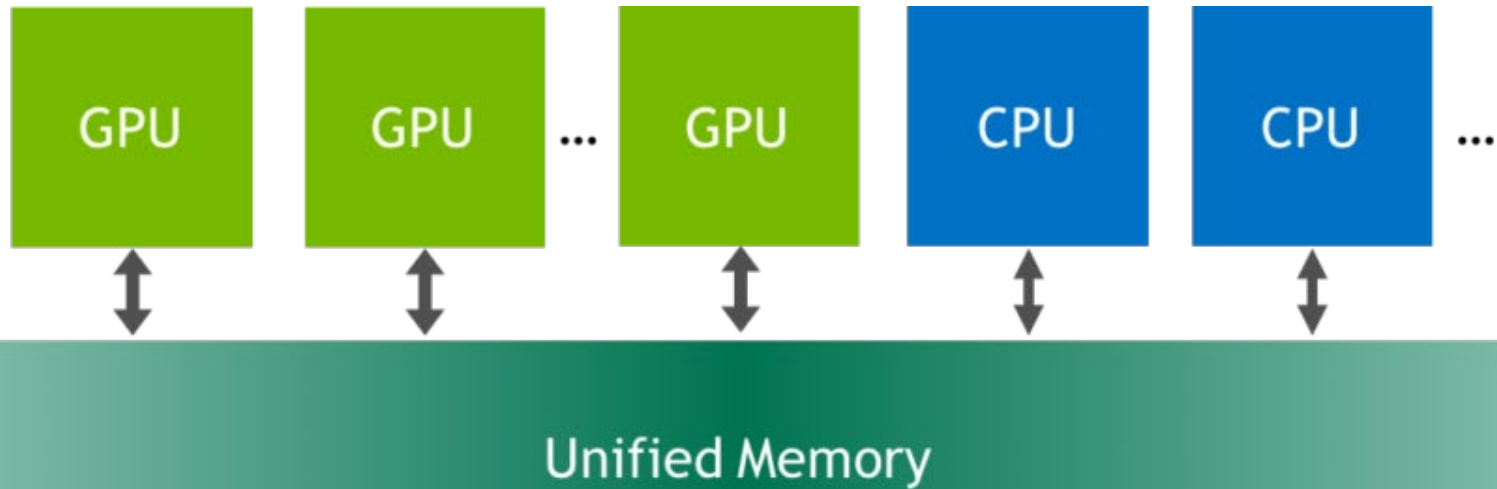
```
__global__ void mykernel(void) {  
}
```

- CUDA C/C++ keyword `__global__` indicates a function that:
  - Runs on the device
  - Is called from host code
- `nvcc` separates source code into host and device components
  - Device functions (e.g. `mykernel()`) processed by NVIDIA compiler
  - Host functions (e.g. `main()`) processed by standard host compiler
    - `gcc, cl.exe`



# Unified Memory

---



```
__host__ __cudaError_t cudaMallocManaged ( void** devPtr, size_t size, unsigned int flags = cudaMemAttachGlobal )
```

Allocates memory that will be automatically managed by the Unified Memory system.

Use the same address pointer in CPU and GPU

# Unified Memory

```
Windows PowerShell
PS C:\ProgramData\NVIDIA Corporation\CUDA Samples\v10.2\bin\win64\Release> .\bandwidthTest.exe
[CUDA Bandwidth Test] - Starting...
Running on...

Device 0: Quadro P1000
Quick Mode

Host to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(GB/s)
  32000000                  13.0

Device to Host Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(GB/s)
  32000000                  12.9

Device to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(GB/s)
  32000000                  82.1

Result = PASS

NOTE: The CUDA Samples are not meant for performance measurements. Results may vary when GPU Boost
is enabled.
PS C:\ProgramData\NVIDIA Corporation\CUDA Samples\v10.2\bin\win64\Release> _
```

# Compute Capability

---



## CUDA-Enabled Jetson Products

### Jetson Products

GPU	Compute Capability
Jetson AGX Xavier	7.2
Jetson Nano	5.3
Jetson TX2	6.2
Jetson TX1	5.3
Tegra X1	5.3

# Compute Capability

<https://en.wikipedia.org/wiki/CUDA>

Technical specifications	Compute capability (version)																	
	1.0	1.1	1.2	1.3	2.x	3.0	3.2	3.5	3.7	5.0	5.2	5.3	6.0	6.1	6.2	7.0 (7.2?)	7.5	
Maximum number of resident grids per device (concurrent kernel execution)	t.b.d.				16		4	32				16	128	32	16	128		
Maximum dimensionality of grid of thread blocks	2				3													
Maximum x-dimension of a grid of thread blocks	65535					2 <sup>31</sup> – 1												
Maximum y-, or z-dimension of a grid of thread blocks	65535																	
Maximum dimensionality of thread block	3																	
Maximum x- or y-dimension of a block	512				1024													
Maximum z-dimension of a block	64																	
Maximum number of threads per block	512				1024													
Warp size	32																	
Maximum number of resident blocks per multiprocessor	8					16					32						16	
Maximum number of resident warps per multiprocessor	24	32		48	64									32				
Maximum number of resident threads per multiprocessor	768	1024		1536	2048										1024			
Number of 32-bit registers per multiprocessor	8 K	16 K		32 K	64 K			128 K	64 K									
Maximum number of 32-bit registers per thread block	N/A			32 K	64 K	32 K	64 K				32 K	64 K		32 K	64 K			
Maximum number of 32-bit registers per thread	124				63		255											
Maximum amount of shared memory per multiprocessor	16 KB				48 KB				112 KB	64 KB	96 KB	64 KB		96 KB	64 KB	96 KB (of 128)	64 KB (of 96)	
Maximum amount of shared memory per thread block	48 KB														48/96 KB	64 KB		
Number of shared memory banks	16				32													
Amount of local memory per thread	16 KB				512 KB													
Constant memory size	64 KB																	
Cache working set per multiprocessor for constant memory	8 KB											4 KB	8 KB					
Cache working set per multiprocessor for texture memory	6 – 8 KB				12 KB		12 – 48 KB			24 KB	48 KB	N/A	24 KB	48 KB	24 KB	32 – 128 KB	32 – 64 KB	
Maximum width for 1D texture reference bound to a CUDA array	8192				65536													
Maximum width for 1D texture reference bound to linear memory	2 <sup>27</sup>																	
Maximum width and number of layers for a 1D layered texture reference	8192 × 512				16384 × 2048													
Maximum width and height for 2D texture reference bound to a CUDA array	65536 × 32768				65536 × 65535													

# CUDA Device Query

```
Windows PowerShell
PS C:\ProgramData\NVIDIA Corporation\CUDA Samples\v10.2\bin\win64\Release> .\deviceQuery.exe
C:\ProgramData\NVIDIA Corporation\CUDA Samples\v10.2\bin\win64\Release\deviceQuery.exe Starting...

  CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)


Device 0: "Quadro P1000"
  CUDA Driver Version / Runtime Version      10.2 / 10.2
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             4096 MBytes (4294967296 bytes)
  ( 4) Multiprocessors, (128) CUDA Cores/MP: 512 CUDA Cores
  GPU Max Clock rate:                       1519 MHz (1.52 GHz)
  Memory Clock rate:                        3004 Mhz
  Memory Bus Width:                         128-bit
  L2 Cache Size:                            524288 bytes
  Maximum Texture Dimension Size (x,y,z)    1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:           65536 bytes
  Total amount of shared memory per block:   49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:       1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                     2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:      Yes with 5 copy engine(s)
  Run time limit on kernels:                 Yes
  Integrated GPU sharing Host Memory:         No
  Support host page-locked memory mapping:   Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                    Disabled
  CUDA Device Driver Mode (TCC or WDDM):      WDDM (Windows Display Driver Model)
  Device supports Unified Addressing (UVA):   Yes
  Device supports Compute Preemption:         Yes
  Supports Cooperative Kernel Launch:        No
  Supports MultiDevice Co-op Kernel Launch:  No
  Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 10.2, CUDA Runtime Version = 10.2, NumDevs = 1
Result = PASS
PS C:\ProgramData\NVIDIA Corporation\CUDA Samples\v10.2\bin\win64\Release> 
```

# Execution Model

---


## Software



Thread




Thread Block




Grid


## Hardware



Scalar Processor



Multiprocessor



Device

Threads are executed by scalar processors

Thread blocks are executed on multiprocessors


Thread blocks do not migrate

Several concurrent thread blocks can reside on one multiprocessor - limited by multiprocessor resources (shared memory and register file)

A kernel is launched as a grid of thread blocks

# Execution Model - Warps

---



A thread block consists of 32-thread warps

A warp is executed physically in parallel (SIMT) on a multiprocessor

SIMT: Single Instruction, Multiple Thread

# Memory Coalescing

---

Global memory access happens in transactions of 32 or 128 bytes

The hardware will try to reduce to as few transactions as possible

*Coalesced* access:


A group of 32 contiguous threads (“warp”) accessing adjacent words

Few transactions and high utilization

*Uncoalesced* access:

A warp of 32 threads accessing scattered words


Many transactions and low utilization





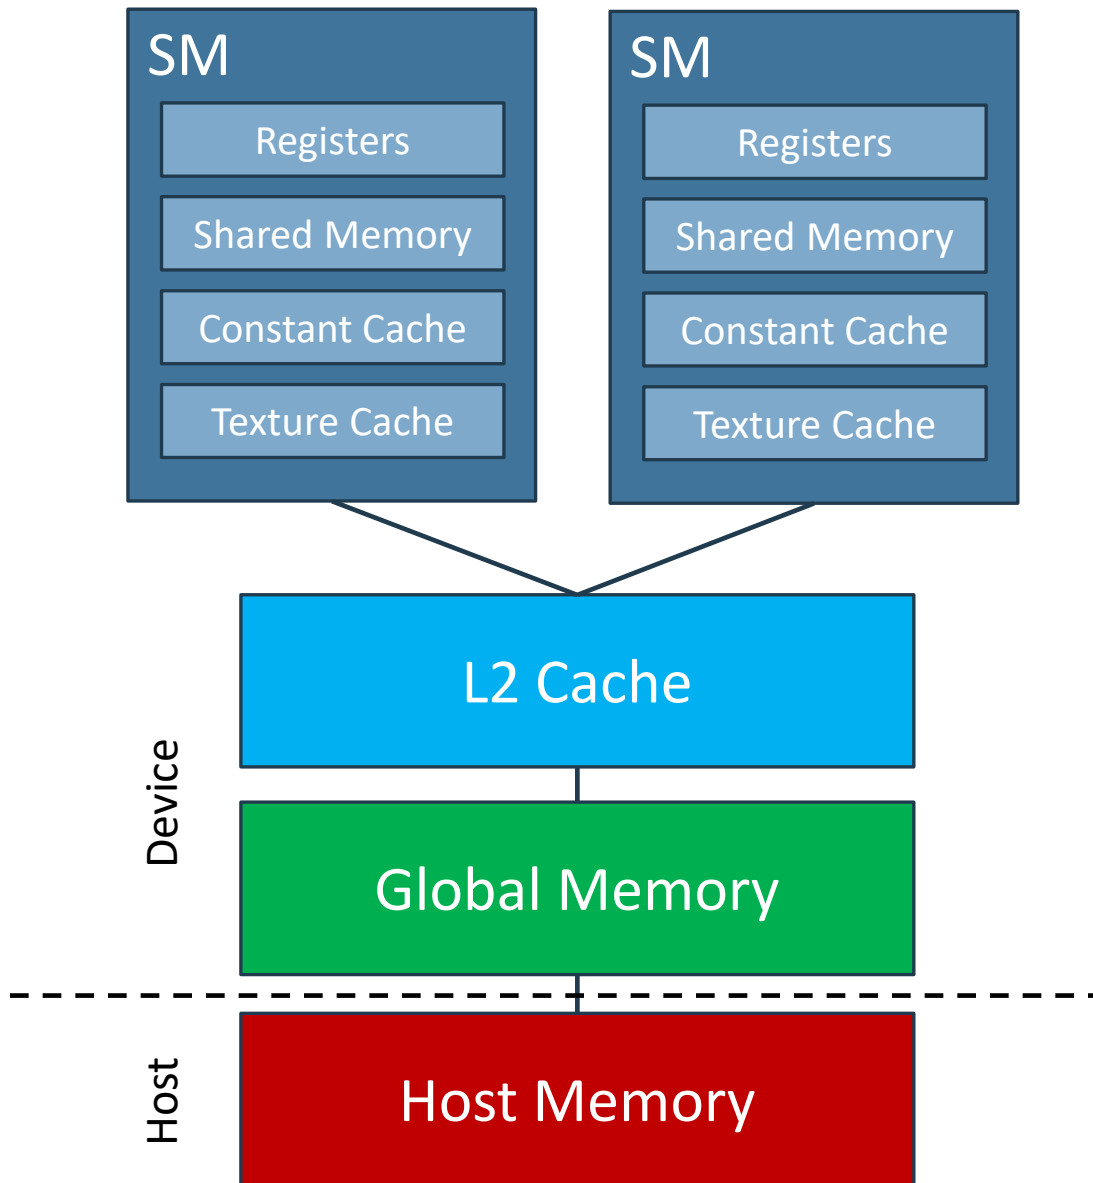
# Warp Divergence

---



50% performance hit!

# Memory Hierarchy



- Global Memory: Large and high latency
- L2 Cache: Medium latency
- SM Caches: Lower latency
- Registers: Lowest latency

# Back to coding again...

---

// Host code

```
MyKernel<<<BLOCKS_PER_GRID, THREADS_PER_BLOCK>>>(...);
```

BLOCKS\_PER\_GRID and THREADS\_PER\_BLOCK are of type *dim3*

Total number of blocks are  $Dg.x * Dg.y * Dg.z$

And total number of threads in the block are  $Db.x * Db.y * Db.z$

*Note: Number of threads per block must be multiple of 32*

```
__global__ void MyKernel(...)  
{  
  
}
```

# Back to coding again...

---

```
// Variable in shared memory
```

```
__shared__ float sum;
```

```
// Device local constant
```

```
__constant__ float growth_rate;
```

```
// Device function to add the elements of two arrays
```

```
__device__ void add(int n, float* x, float* y)
```

```
{
```

```
    for (int i = 0; i < n; i++)
```

```
        y[i] = x[i] + y[i];
```

```
}
```

```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    for (int i = 0; i < n; i++)
        y[i] = x[i] + y[i];
}

int main(void)
{
    int N = 1 << 20;
    float* x, * y;

    // Allocate Unified Memory - accessible from CPU or GPU
    cudaMallocManaged(&x, N * sizeof(float));
    cudaMallocManaged(&y, N * sizeof(float));

    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }

    // Run kernel on 1M elements on the GPU
    add << <1, 1 >> > (N, x, y);

    // Wait for GPU to finish before accessing on host
    cudaDeviceSynchronize();

    // Check for errors (all values should be 3.0f)
    float maxError = 0.0f;
    for (int i = 0; i < N; i++)
        maxError = fmax(maxError, fabs(y[i] - 3.0f));
    std::cout << "Max error: " << maxError << std::endl;

    // Free memory
    cudaFree(x);
    cudaFree(y);

    return 0;
}

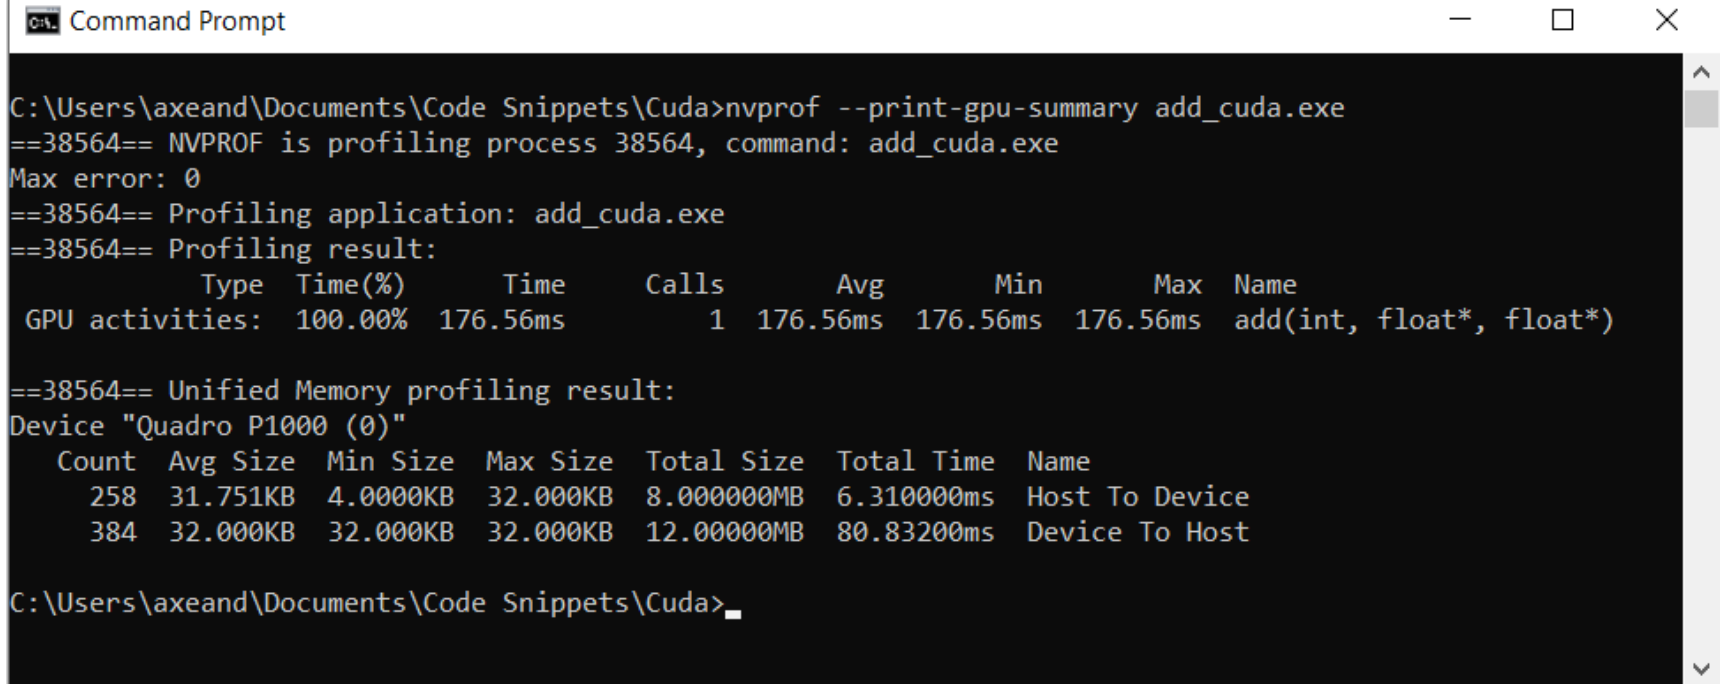
```

```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    for (int i = 0; i < n; i++)
        y[i] = x[i] + y[i];
}

int main(void)
{

```



```

C:\Users\axeand\Documents\Code Snippets\Cuda>nvprof --print-gpu-summary add_cuda.exe
==38564== NVPROF is profiling process 38564, command: add_cuda.exe
Max error: 0
==38564== Profiling application: add_cuda.exe
==38564== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max  Name
GPU activities: 100.00%   176.56ms        1  176.56ms  176.56ms  176.56ms  add(int, float*, float*)

==38564== Unified Memory profiling result:
Device "Quadro P1000 (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
     258   31.751KB  4.0000KB  32.000KB   8.000000MB   6.310000ms  Host To Device
     384   32.000KB  32.000KB  32.000KB  12.000000MB   80.83200ms  Device To Host

C:\Users\axeand\Documents\Code Snippets\Cuda>

```

```

float maxError = 0.0f;
for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i] - 3.0f));
std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x);
cudaFree(y);

return 0;
}

```

```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    int index = threadIdx.x;
    int stride = blockDim.x;
    for (int i = index; i < n; i += stride)
        y[i] = x[i] + y[i];
}

int main(void)
{
    int N = 1 << 20;
    float* x, * y;

    // Allocate Unified Memory - accessible from CPU or GPU
    cudaMallocManaged(&x, N * sizeof(float));
    cudaMallocManaged(&y, N * sizeof(float));

    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }

    // Run kernel on 1M elements on the GPU
    add << <1, 256 >> > (N, x, y);

    // Wait for GPU to finish before accessing on host
    cudaDeviceSynchronize();

    // Check for errors (all values should be 3.0f)
    float maxError = 0.0f;
    for (int i = 0; i < N; i++)
        maxError = fmax(maxError, fabs(y[i] - 3.0f));
    std::cout << "Max error: " << maxError << std::endl;

    // Free memory
    cudaFree(x);
    cudaFree(y);

    return 0;
}

```

```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    int index = threadIdx.x;
    int stride = blockDim.x;
    for (int i = index; i < n; i += stride)
        y[i] = x[i] + y[i];
}

```

```

C:\Users\axeand\Documents\Code Snippets\Cuda>nvprof --print-gpu-summary add_cuda.exe
==1716== NVPROF is profiling process 1716, command: add_cuda.exe
Max error: 0
==1716== Profiling application: add_cuda.exe
==1716== Profiling result:
   Type  Time(%)      Time   Calls    Avg      Min      Max  Name
GPU activities: 100.00%  1.2942ms        1  1.2942ms  1.2942ms  1.2942ms  add(int, float*, float*)

==1716== Unified Memory profiling result:
Device "Quadro P1000 (0)"
  Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    258   31.751KB  4.0000KB  32.000KB   8.000000MB   6.331800ms  Host To Device
    384   32.000KB  32.000KB  32.000KB  12.000000MB  81.23890ms  Device To Host

C:\Users\axeand\Documents\Code Snippets\Cuda>

```

```

// Error handling
float maxError = 0.0f;
for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i] - 3.0f));
std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x);
cudaFree(y);

return 0;
}

```



```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    int index = blockIdx.x * blockDim.x + threadIdx.x;
    int stride = blockDim.x * gridDim.x;
    for (int i = index; i < n; i += stride)
        y[i] = x[i] + y[i];
}

int main(void)
{
    int N = 1 << 20;
    float* x, * y;

    // Allocate Unified Memory - accessible from CPU or GPU
    cudaMallocManaged(&x, N * sizeof(float));
    cudaMallocManaged(&y, N * sizeof(float));

    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }

    int blockSize = 256;
    int numBlocks = (N + blockSize - 1) / blockSize;

    // Run kernel on 1M elements on the GPU
    add << <numBlocks, blockSize >> > (N, x, y);

    // Wait for GPU to finish before accessing on host
    cudaDeviceSynchronize();

    // Check for errors (all values should be 3.0f)
    float maxError = 0.0f;
    for (int i = 0; i < N; i++)
        maxError = fmax(maxError, fabs(y[i] - 3.0f));
    std::cout << "Max error: " << maxError << std::endl;

    // Free memory
    cudaFree(x);
    cudaFree(y);

    return 0;
}

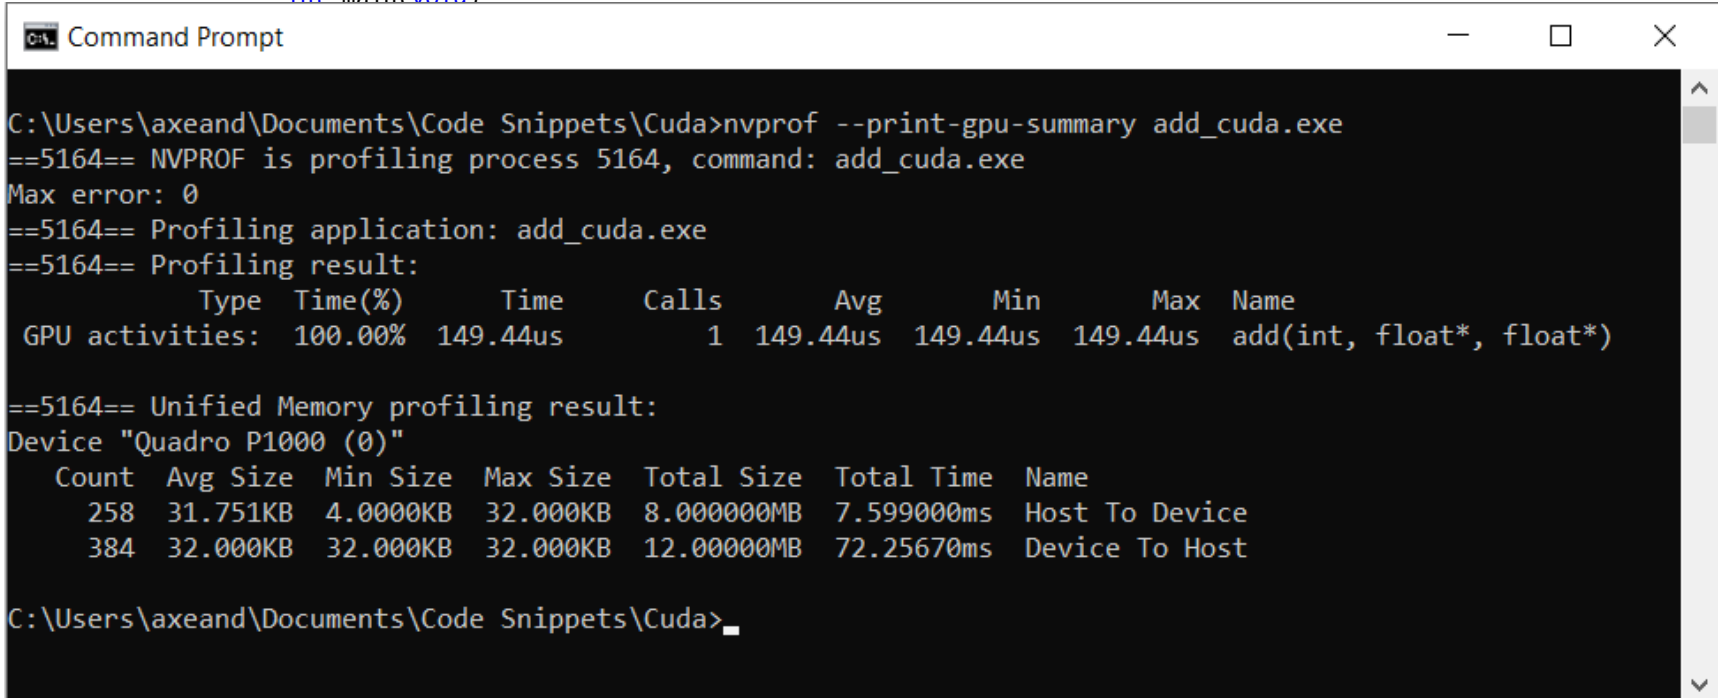
```

```

#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__ void add(int n, float* x, float* y)
{
    int index = blockIdx.x * blockDim.x + threadIdx.x;
    int stride = blockDim.x * gridDim.x;
    for (int i = index; i < n; i += stride)
        y[i] = x[i] + y[i];
}

int main(void)

```



```

C:\Users\axeand\Documents\Code Snippets\Cuda>nvprof --print-gpu-summary add_cuda.exe
==5164== NVPROF is profiling process 5164, command: add_cuda.exe
Max error: 0
==5164== Profiling application: add_cuda.exe
==5164== Profiling result:
          Type  Time(%)      Time   Calls    Avg      Min      Max  Name
GPU activities: 100.00%   149.44us        1  149.44us  149.44us  149.44us  add(int, float*, float*)

==5164== Unified Memory profiling result:
Device "Quadro P1000 (0)"
  Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    258   31.751KB  4.0000KB  32.000KB   8.000000MB   7.599000ms  Host To Device
    384   32.000KB  32.000KB  32.000KB  12.000000MB  72.25670ms  Device To Host

C:\Users\axeand\Documents\Code Snippets\Cuda>_

```

```

// Check for errors (all values should be 3.0f)
float maxError = 0.0f;
for (int i = 0; i < N; i++)
    maxError = fmax(maxError, fabs(y[i] - 3.0f));
std::cout << "Max error: " << maxError << std::endl;

// Free memory
cudaFree(x);
cudaFree(y);

return 0;
}

```

# Some tips for the road...

---

- In the kernel invocation, `<<<Blocks, Threads>>>`, try to choose a number of threads that divides evenly with the number of threads in a warp. If you don't, you end up with launching a block that contains inactive threads.
- In your kernel, try to have each thread in a warp follow the same code path. If you don't, you get what's called warp divergence. This happens because the GPU has to run the entire warp through each of the divergent code paths.
- In your kernel, try to have each thread in a warp load and store data in specific patterns. For instance, have the threads in a warp access consecutive 32-bit words in global memory.

# CUDA in Python - PyCUDA

---

```
import pycuda.compiler as comp
import pycuda.driver as drv
import numpy
import pycuda.autoinit

mod = comp.SourceModule("""
__global__ void multiply_them(float *dest, float *a, float *b)
{
    const int i = threadIdx.x;
    dest[i] = a[i] * b[i];
}
""")

multiply_them = mod.get_function("multiply_them")

a = numpy.random.randn(400).astype(numpy.float32)
b = numpy.random.randn(400).astype(numpy.float32)

dest = numpy.zeros_like(a)
multiply_them(
    drv.Out(dest), drv.In(a), drv.In(b),
    block=(400,1,1))

print dest-a*b
```

<https://pypi.org/project/pycuda/>

# CUDA Samples

cuda-samples/Samples at master · NVIDIA · GitHub

github.com/NVIDIA/cuda-samples/tree/master/Samples

Apps Embedded Progra... JU e-meeting - Zoom Launch Meeting - Z... DH-Self Service


NVIDIA / cuda-samples

Code Issues 6 Pull requests 4 Actions Projects 0 Wiki S

Branch: master cuda-samples / Samples /

mdoijade Add and update samples for cuda 10.2 support Latest commit 6be5146 on Oct 23, 2019

..		
EGLStream_CUDA_Interop	Add and update samples for cuda 10.2 support	7 months ago
MersenneTwisterGP11213	Add and update samples for cuda 10.2 support	7 months ago
NV12toBGRandResize	Add and update samples for cuda 10.2 support	7 months ago
UnifiedMemoryPerf	Add and update samples for cuda 10.2 support	7 months ago
bandwidthTest	Add and update samples for cuda 10.2 support	7 months ago
boxFilterNPP	Add and update samples for cuda 10.2 support	7 months ago
cannyEdgeDetectorNPP	Add and update samples for cuda 10.2 support	7 months ago
conjugateGradientCudaGraphs	Add and update samples for cuda 10.2 support	7 months ago
conjugateGradientMultiBlockCG	Add and update samples for cuda 10.2 support	7 months ago
conjugateGradientMultiDeviceCG	Add and update samples for cuda 10.2 support	7 months ago
cuSolverDn_LinearSolver	Add and update samples for cuda 10.2 support	7 months ago



Create new file Upload files Find file History

# Introduction to CUDA

---

## Questions?

Contact information

Andreas Axelsson

Email: [andreas.axelsson@ju.se](mailto:andreas.axelsson@ju.se)

Mobile: 0709-467760