

Contemporary Computer Architecture TDSN13

LECTURE 7 – CUDA MEMORY ACCESS PATTERNS


ANDREAS AXELSSON (ANDREAS.AXELSSON@JU.SE)

DRAM Burst – A System View




- Each address space is partitioned into burst sections
 - Whenever a location is accessed, all other locations in the same section are also delivered to the processor
- Basic example: a 16-byte address space, 4-byte burst sections
 - In practice, we have at least 4GB address space, burst section sizes of 128-bytes or more

Memory Coalescing



- When all threads of a warp execute a load instruction, if all accessed locations fall into the same burst section, only one DRAM request will be made and the access is fully coalesced.

Memory Coalescing

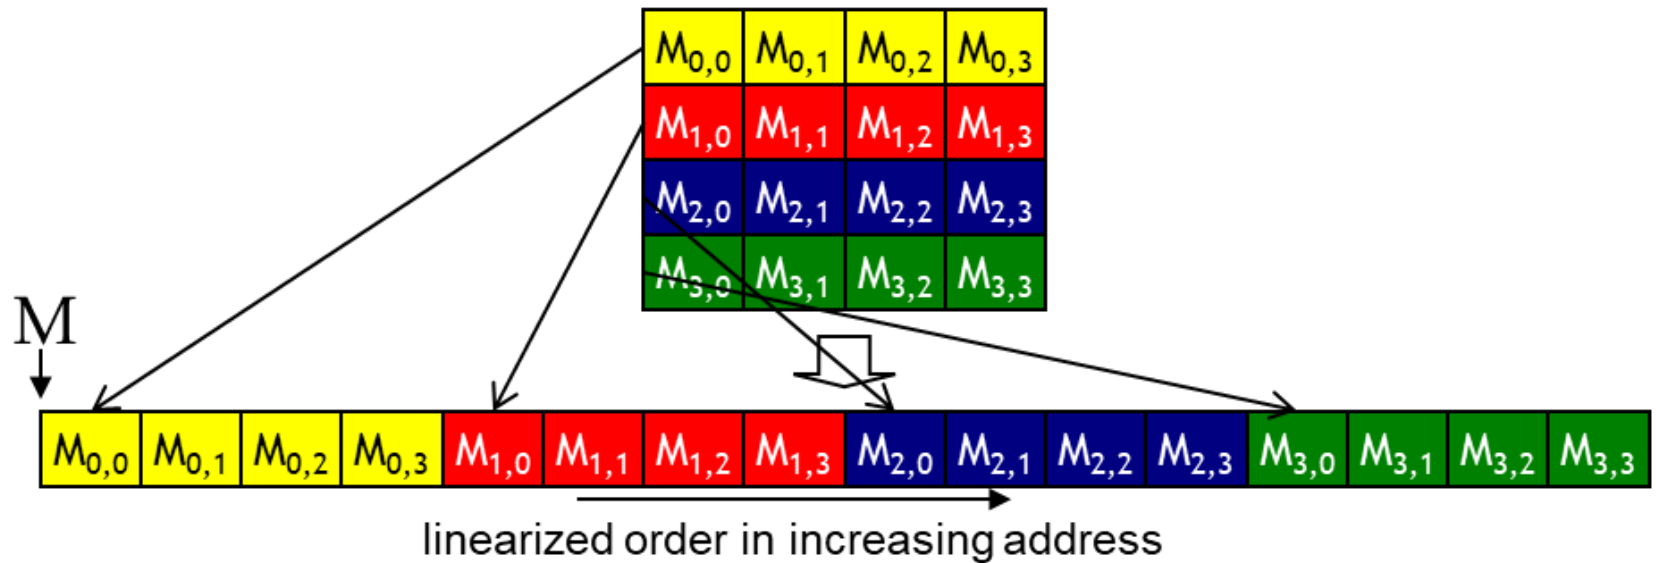


- When the accessed locations spread across burst section boundaries:
 - Coalescing fails
 - Multiple DRAM requests are made
 - The access is not fully coalesced.
- Some of the bytes accessed and transferred are not used by the threads


Is access coalesced?

- Accesses in a warp are to consecutive locations if the index in an array access is in the form of
 - $A[(\text{expression with terms independent of threadIdx.x}) + \text{threadIdx.x}]$;

2D Array in Linear Memory Space



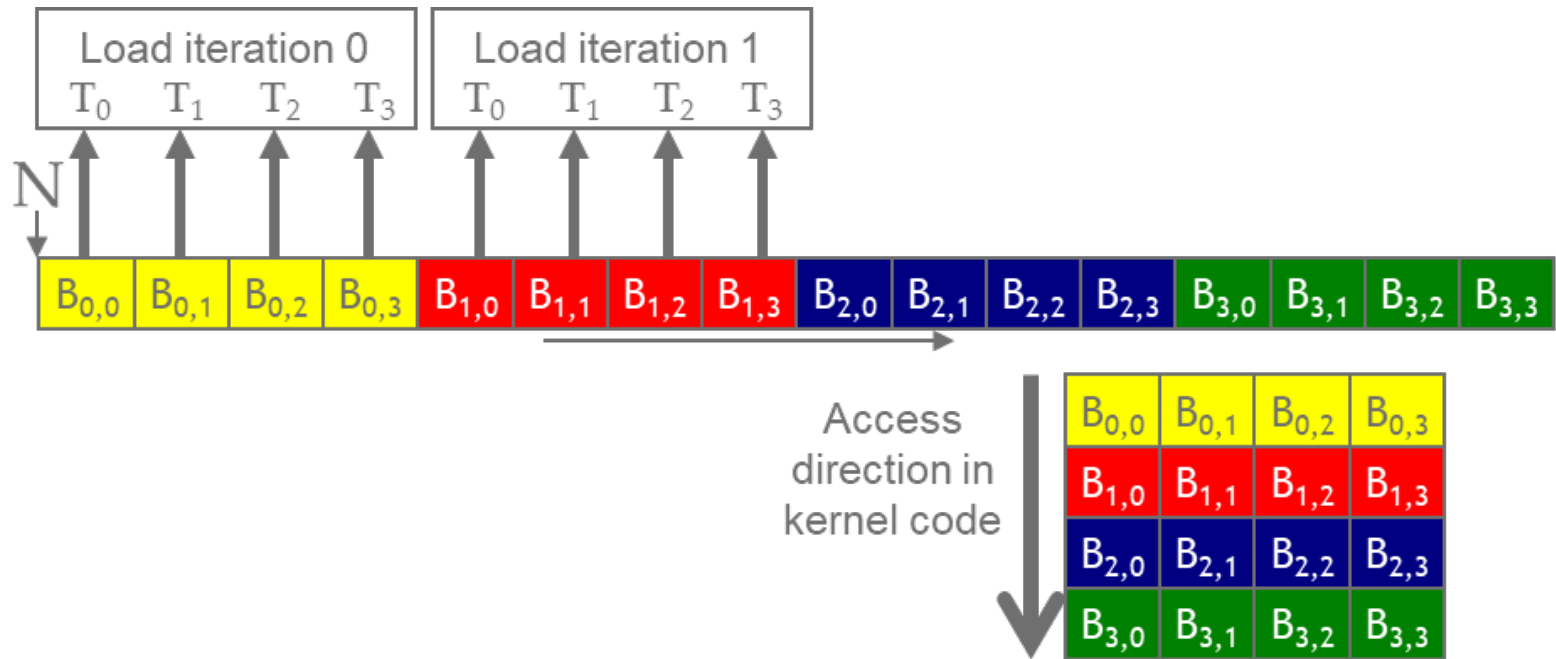
Access Pattern Multiplication



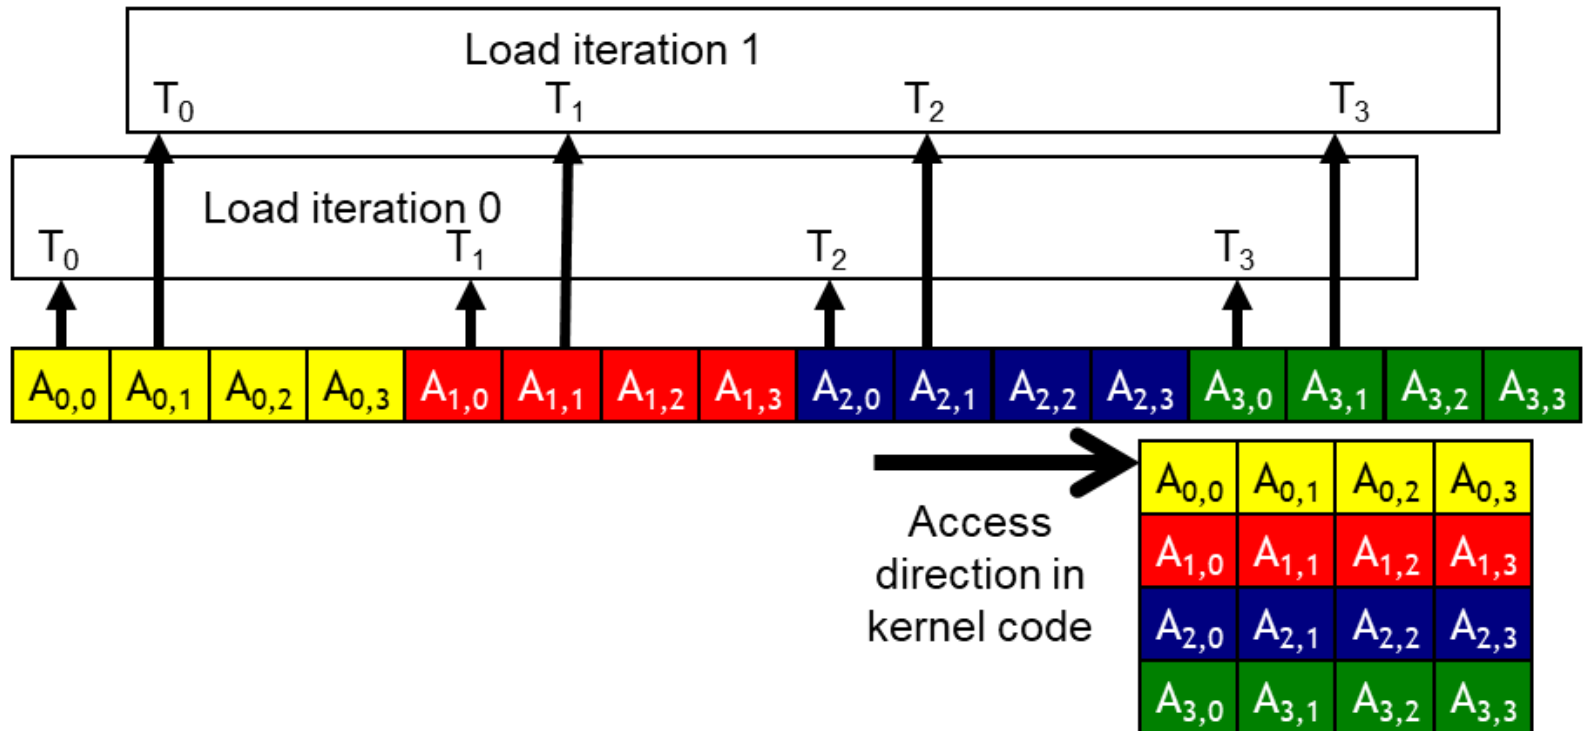
i is the loop counter in the inner product loop of the kernel code

A is $m \times n$, B is $n \times k$
 $\text{Col} = \text{blockIdx.x} \times \text{blockDim.x} + \text{threadIdx.x}$

B accesses are coalesced



A accesses are not coalesced



- Coalesced accesses in a warp are to consecutive locations if the index in an array access is in the form of
 - $A[(\text{expression with terms independent of threadIdx.x}) + \text{threadIdx.x}]$;

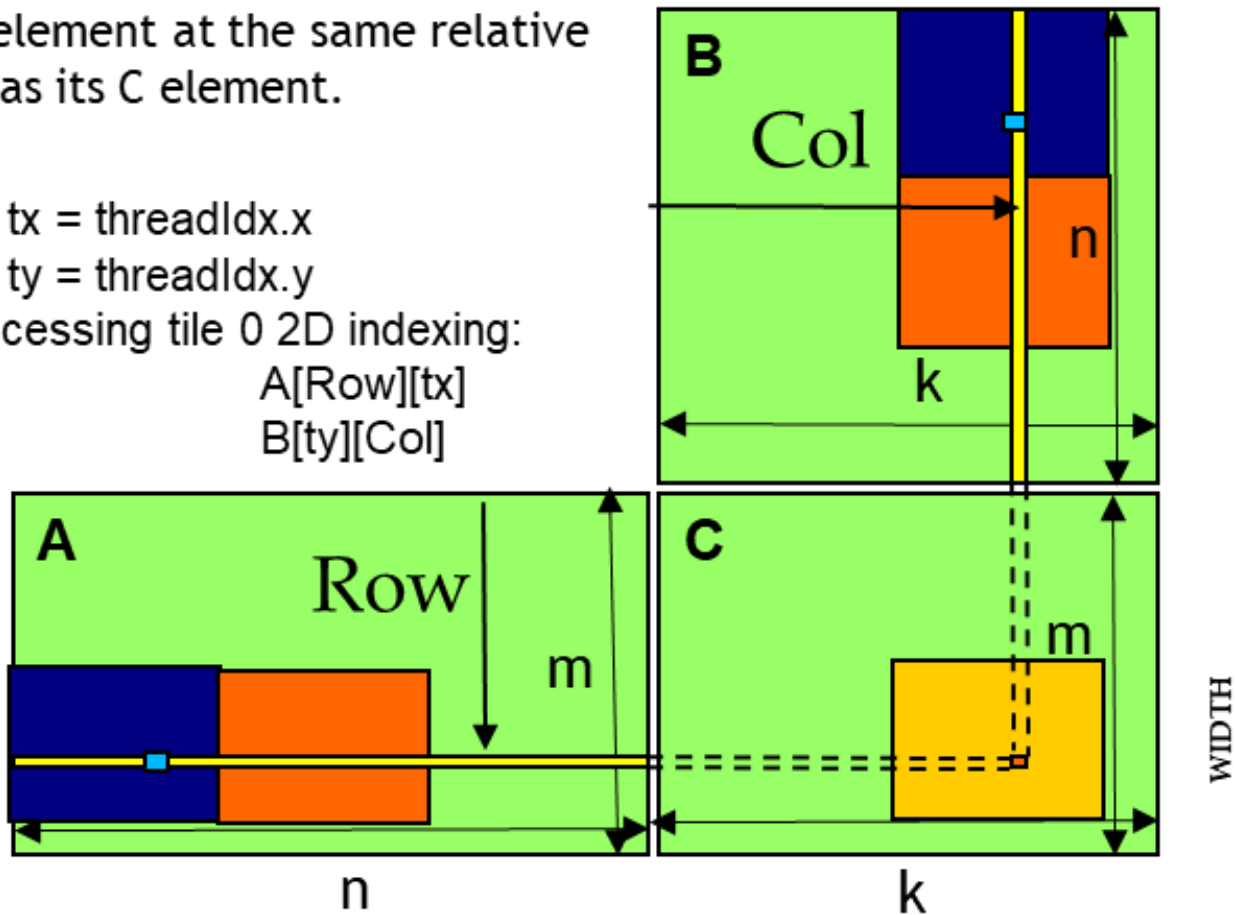
Corner Turning

- What if we must iterate through data along the row direction?
 - We can use the shared memory to enable memory coalescing — this is the technique called Corner Turning
- how?
 - Use a tiled algorithm
 - Threads of a block can first cooperatively load the tiles into the shared memory
 - Care must be taken to ensure that these tiles are loaded in a coalesced pattern
 - Once the data is in shared memory, they can be accessed either on a row basis or a column basis with much less performance variation because the shared memories are implemented as intrinsically high-speed on-chip memory that does not require coalescing to achieve high data access rate.

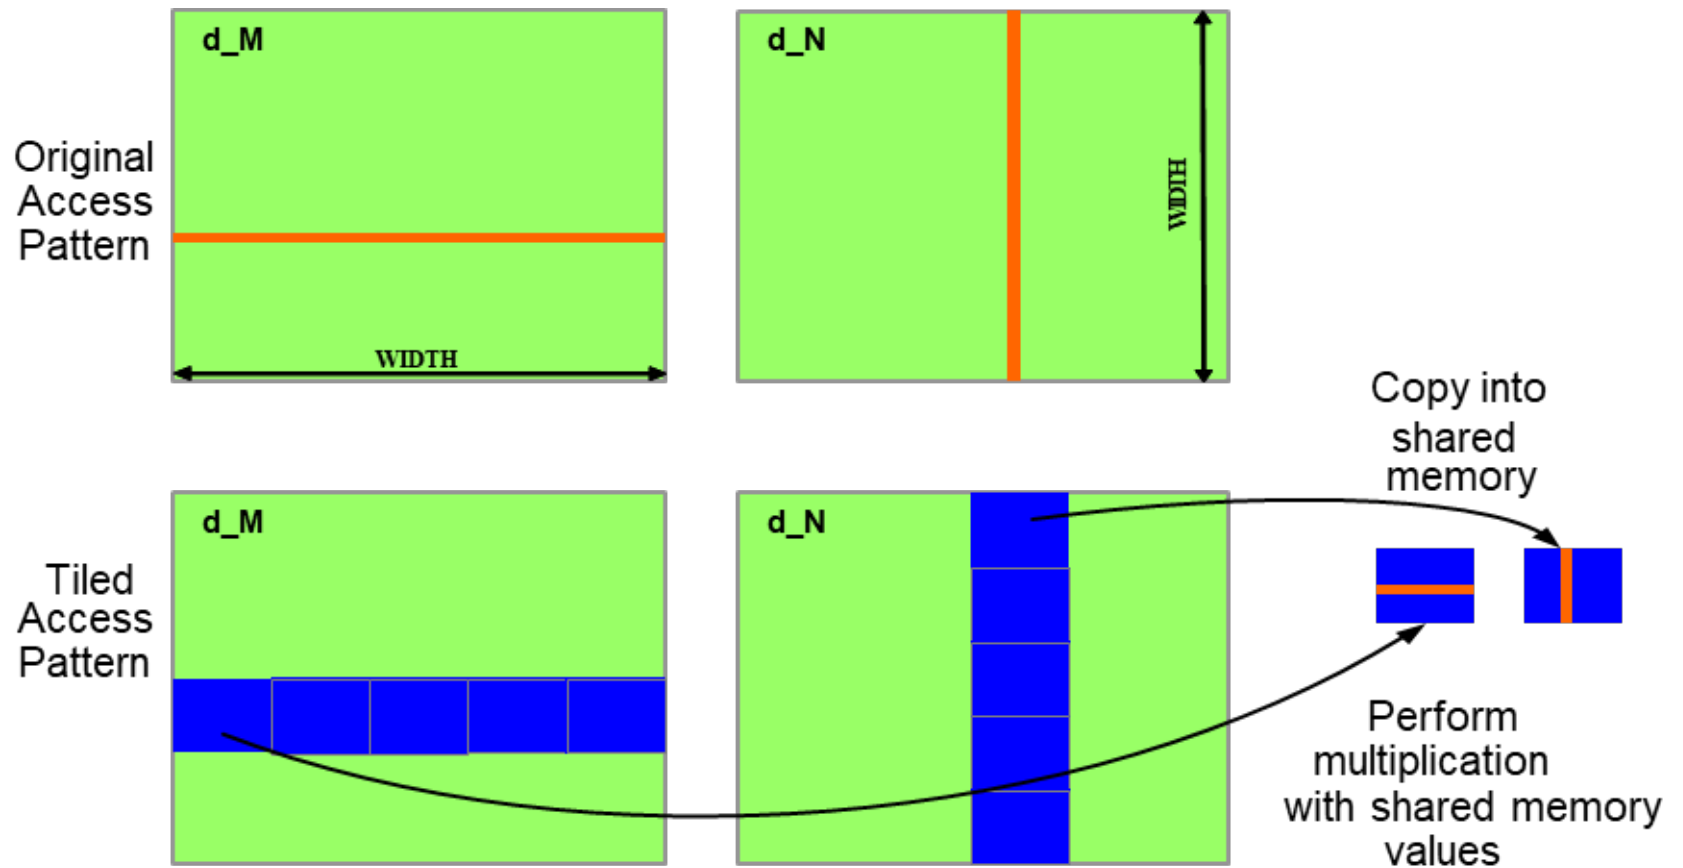
Corner Turning – Loading a tile

Have each thread load an A element and a B element at the same relative position as its C element.

```
int tx = threadIdx.x  
int ty = threadIdx.y  
Accessing tile 0 2D indexing:  
A[Row][tx]  
B[ty][Col]
```



Corner Turning



Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] *
ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

Remember that this Kernel
runs $TILE_WIDTH \times TILE_WIDTH$
number of threads in parallel
in each block (SM)

Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();

        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] *
            ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, Int Width)
{
    __shared__ float ds_M[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_N[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;  int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int Row = by * blockDim.y + ty;
    int Col = bx * blockDim.x + tx;
    float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    for (int p = 0; p < n/TILE_WIDTH; ++p) {
        // Collaborative loading of M and N tiles into shared memory
        ds_M[ty][tx] = M[Row*Width + p*TILE_WIDTH+tx];
        ds_N[ty][tx] = N[(t*TILE_WIDTH+ty)*Width + Col];
        __syncthreads();


        for (int i = 0; i < TILE_WIDTH; ++i) Pvalue += ds_M[ty][i] *
ds_N[i][tx];
        __syncthreads();
    }
    P[Row*Width+Col] = Pvalue;
}
```

Histogram

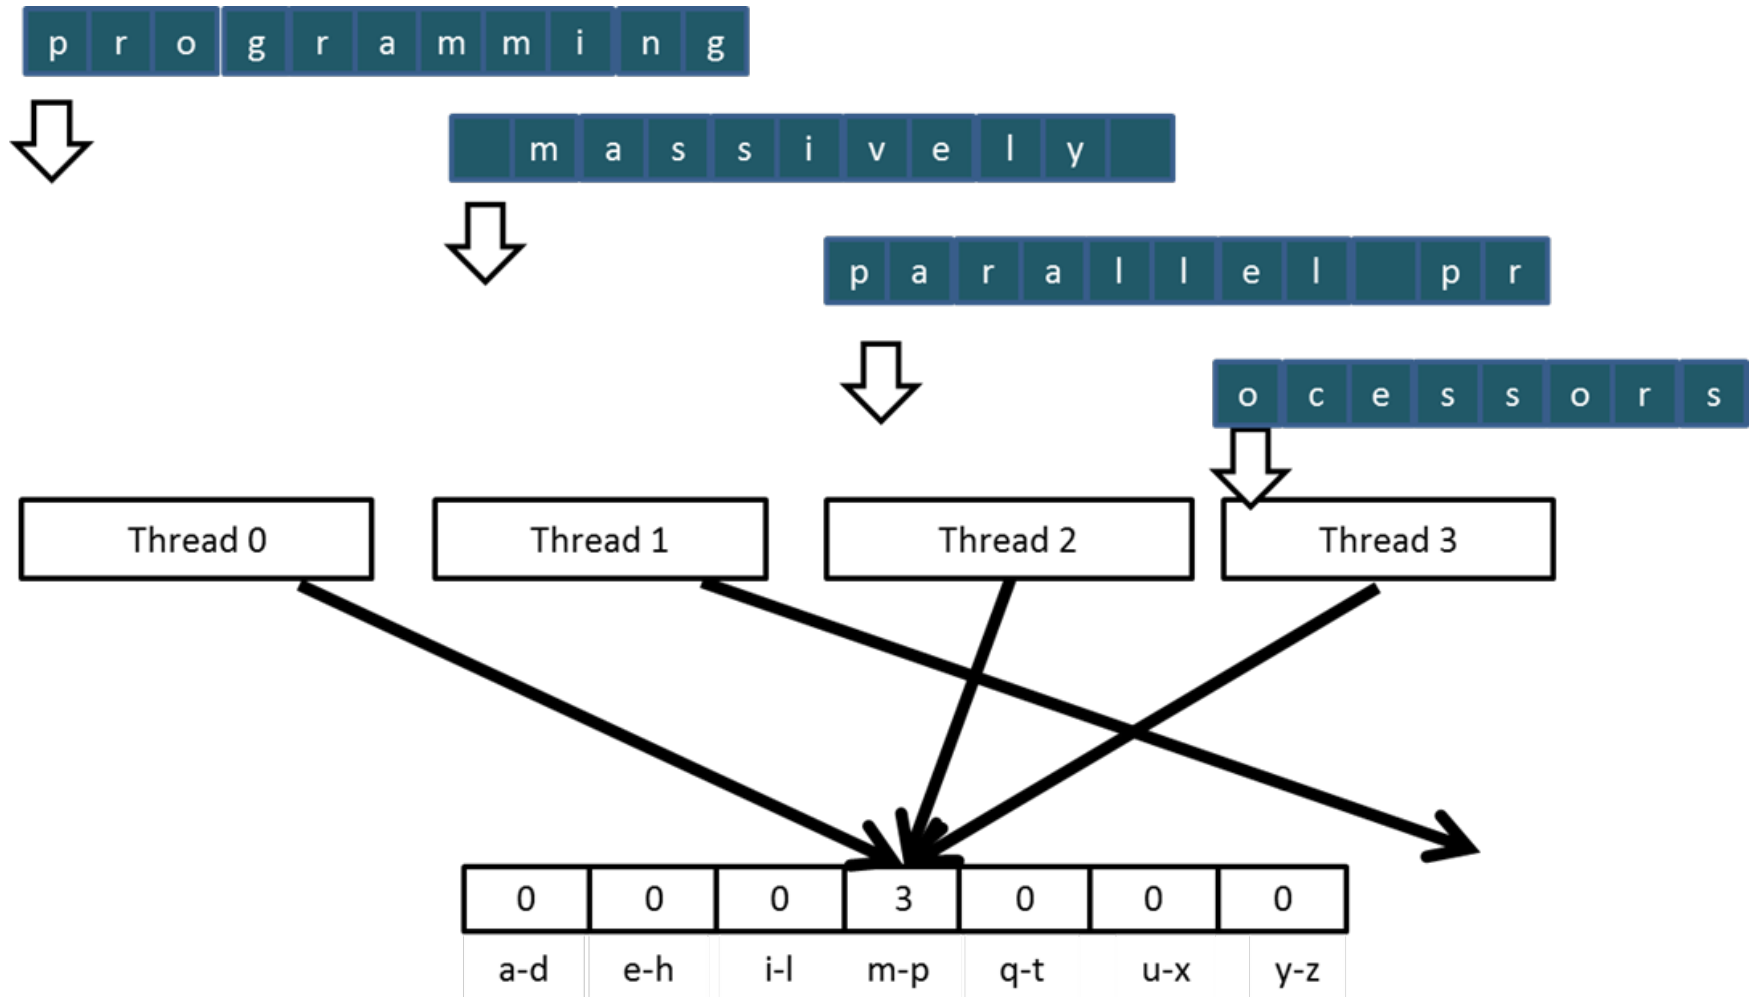
- A method for extracting notable features and patterns from large data sets
 - Feature extraction for object recognition in images
 - Fraud detection in credit card transactions
 - Correlating heavenly object movements in astrophysics
 - ...
- Basic histograms - for each element in the data set, use the value to identify a “bin counter” to increment

A Text Histogram Example

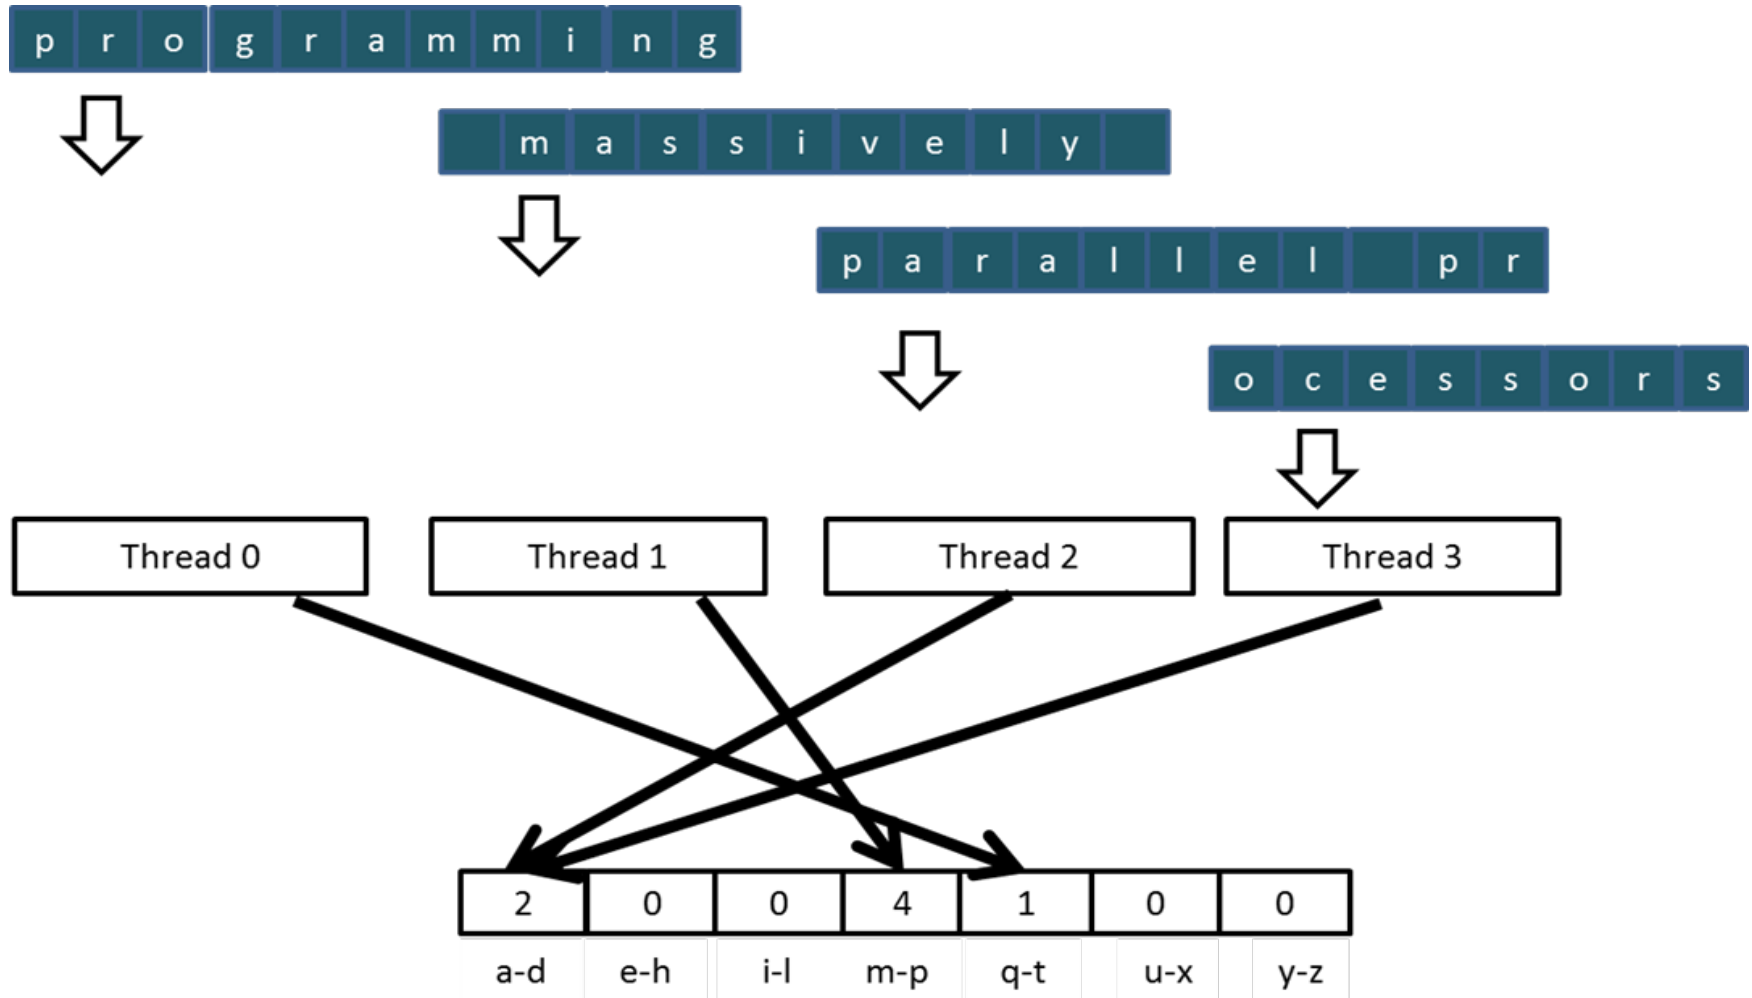
- Define the bins as four-letter sections of the alphabet: a-d, e-h, i-l, n-p, ...
- For each character in an input string, increment the appropriate bin counter.
- In the phrase “Programming Massively Parallel Processors” the output histogram is shown below:



Sectioned Partition (Iteration #1)




Sectioned Partition (Iteration #2)



Partitioning Access Efficiency

- Sectioned partitioning results in poor memory access efficiency
 - Adjacent threads do not access adjacent memory locations
 - Accesses are not coalesced
 - DRAM bandwidth is poorly utilized

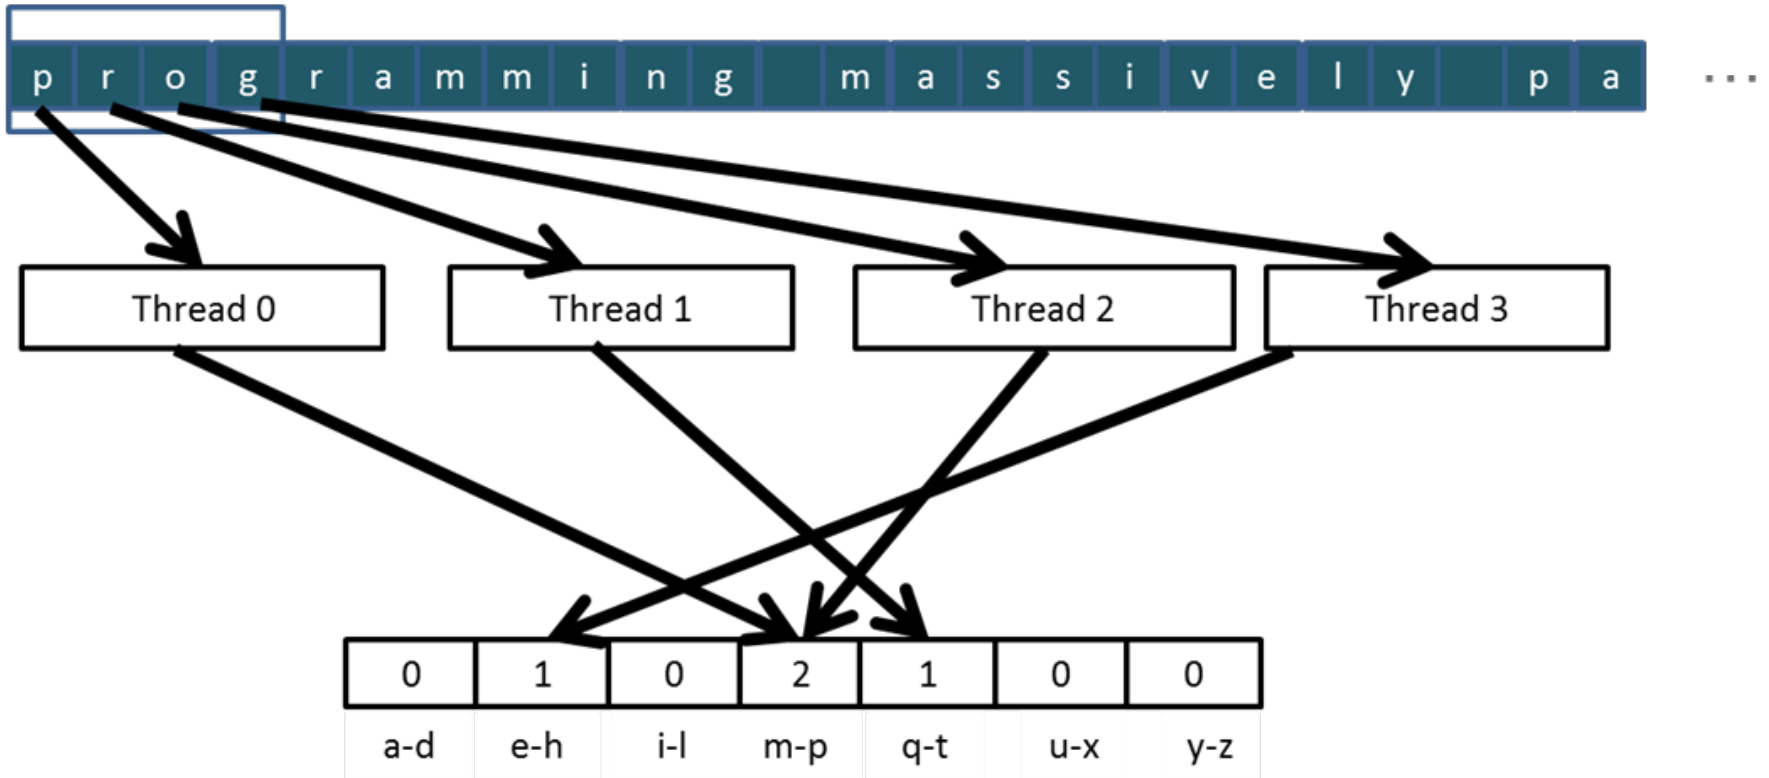


- Change to interleaved partitioning
 - All threads process a contiguous section of elements
 - They all move to the next section and repeat
 - The memory accesses are coalesced

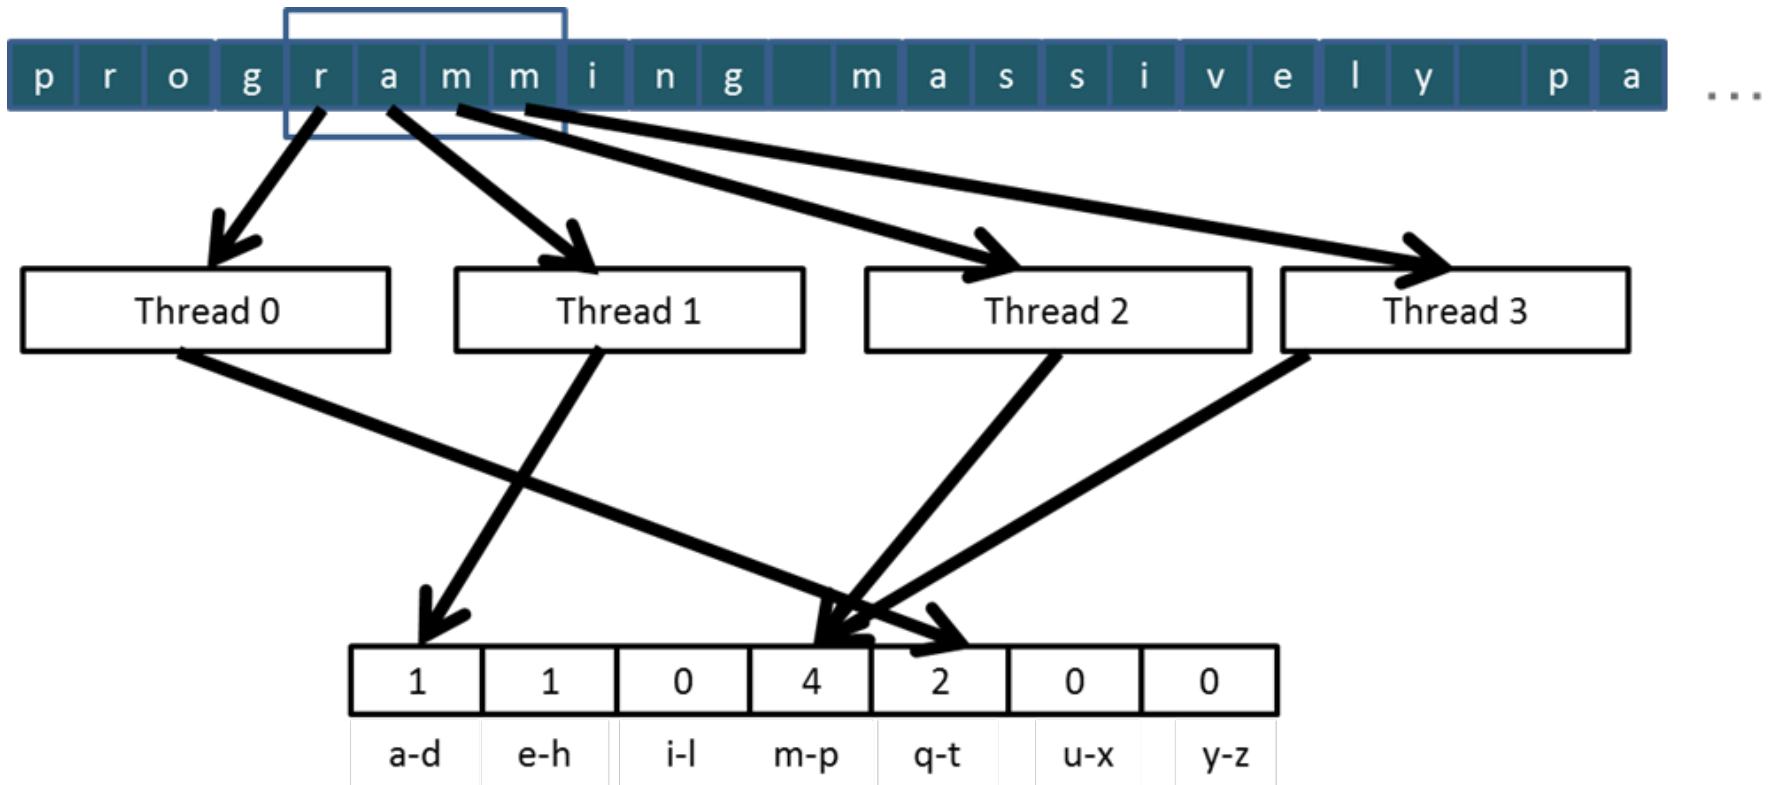


Interleaved Partitioning of Input

- For coalescing and better memory access performance



Interleaved Partitioning (Iter #2)



Data Race in Parallel Threads

thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

Old and New are per-thread register variables.

Question 1: If Mem[x] was initially 0, what would the value of Mem[x] be after threads 1 and 2 have completed?

Question 2: What does each thread get in their Old variable?

Unfortunately, the answers may vary according to the relative execution timing between the two threads, which is referred to as a **data race**.

Timing Scenario #1

| Time | Thread 1 | Thread 2 |
|------|--|--|
| 1 | (0) $\text{Old} \leftarrow \text{Mem}[x]$ | |
| 2 | (1) $\text{New} \leftarrow \text{Old} + 1$ | |
| 3 | (1) $\text{Mem}[x] \leftarrow \text{New}$ | |
| 4 | | (1) $\text{Old} \leftarrow \text{Mem}[x]$ |
| 5 | | (2) $\text{New} \leftarrow \text{Old} + 1$ |
| 6 | | (2) $\text{Mem}[x] \leftarrow \text{New}$ |

- Thread 1 Old = 0
- Thread 2 Old = 1
- $\text{Mem}[x] = 2$ after the sequence

Timing Scenario #2

| Time | Thread 1 | Thread 2 |
|------|------------------------------|------------------------------|
| 1 | | (0) Old \leftarrow Mem[x] |
| 2 | | (1) New \leftarrow Old + 1 |
| 3 | | (1) Mem[x] \leftarrow New |
| 4 | (1) Old \leftarrow Mem[x] | |
| 5 | (2) New \leftarrow Old + 1 | |
| 6 | (2) Mem[x] \leftarrow New | |

- Thread 1 Old = 1
- Thread 2 Old = 0
- Mem[x] = 2 after the sequence

Timing Scenario #3

| Time | Thread 1 | Thread 2 |
|------|------------------------------|------------------------------|
| 1 | (0) Old \leftarrow Mem[x] | |
| 2 | (1) New \leftarrow Old + 1 | |
| 3 | | (0) Old \leftarrow Mem[x] |
| 4 | (1) Mem[x] \leftarrow New | |
| 5 | | (1) New \leftarrow Old + 1 |
| 6 | | (1) Mem[x] \leftarrow New |

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Timing Scenario #4

| Time | Thread 1 | Thread 2 |
|------|------------------------------|------------------------------|
| 1 | | (0) Old \leftarrow Mem[x] |
| 2 | | (1) New \leftarrow Old + 1 |
| 3 | (0) Old \leftarrow Mem[x] | |
| 4 | | (1) Mem[x] \leftarrow New |
| 5 | (1) New \leftarrow Old + 1 | |
| 6 | (1) Mem[x] \leftarrow New | |

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Key Concept of Atomic Operations

- A read-modify-write operation performed by a single hardware instruction on a memory location *address*
 - Read the old value, calculate a new value, and write the new value to the location
- The hardware ensures that no other threads can perform another read-modify-write operation on the same location until the current atomic operation is complete
 - Any other threads that attempt to perform an atomic operation on the same location will typically be held in a queue
 - All threads perform their atomic operations **serially** on the same location

Atomic Operations in CUDA

- Performed by calling functions that are translated into single instructions (a.k.a. *intrinsic functions* or *intrinsics*)
 - Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)
 - Read CUDA C programming Guide for details
- Atomic Add
 - ```
int atomicAdd(int* address, int val);
```
  - reads the 32-bit word **old** from the location pointed to by **address** in global or shared memory, computes (**old + val**), and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns **old**.

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

# More Atomic Adds in CUDA

---

- Unsigned 32-bit integer atomic add

```
unsigned int atomicAdd(unsigned int* address,
 unsigned int val);
```

- Unsigned 64-bit integer atomic add

```
unsigned long long int atomicAdd(unsigned long long
 int* address, unsigned long long int val);
```

- Single-precision floating-point atomic add (Compute capability 2.x+)

```
float atomicAdd(float* address, float val);
```

- Double-precision floating-point atomic add (Compute capability 6.x+)

```
double atomicAdd(double* address, double val);
```

- 16-bit floating-point atomic add (Compute capability 7.x+)

```
__half atomicAdd(__half* address, __half val);
```

# A Generic Text Histogram Kernel

---

- The kernel receives a pointer to the input buffer of byte values
- Each thread process the input in a strided pattern

```
__global__ void histo_kernel(unsigned char *buffer,
 long size, unsigned int *histo)
{
 int i = threadIdx.x + blockIdx.x * blockDim.x;

 // stride is total number of threads
 int stride = blockDim.x * gridDim.x;

 // All threads handle blockDim.x * gridDim.x
 // consecutive elements
 while (i < size) {
 atomicAdd(&(histo[buffer[i]]), 1);
 i += stride;
 }
}
```

# Our Text Histogram Kernel (cont.)

---

- The kernel receives a pointer to the input buffer of byte values
- Each thread process the input in a strided pattern

```
__global__ void histo_kernel(unsigned char *buffer,
 long size, unsigned int *histo)
{
 int i = threadIdx.x + blockIdx.x * blockDim.x;

 // stride is total number of threads
 int stride = blockDim.x * gridDim.x;

 // All threads handle blockDim.x * gridDim.x
 // consecutive elements
 while (i < size) {
 int alphabet_position = buffer[i] - "a";
 if (alphabet_position >= 0 && alphabet_position < 26)
 atomicAdd(&(histo[alphabet_position/4]), 1);
 i += stride;
 }
}
```



# Introduction to CUDA

---

## Questions?

Contact information

Andreas Axelsson

Email: [andreas.axelsson@ju.se](mailto:andreas.axelsson@ju.se)

Mobile: 0709-467760